



# Are we there yet?

Lessons Learnt Performance Testing A Crypto Exchange API

**Ben Shi**

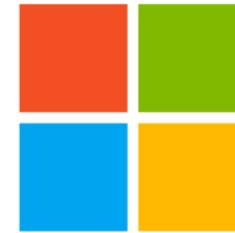
#dddsydney

[@hbishi](https://twitter.com/hbishi) | [/in/benshi](https://www.linkedin.com/in/benshi) | [hbishi.com](https://hbishi.com)



👋 hi

**GitHub**



**Microsoft**



**Telstra  
Purple**

**NDC { Sydney }**

**YOW!**



Backed by Thinkmill



# About Me



Quality Assurance Engineer

~3.5 years

J.P.Morgan



Software Engineer / Consultant

~7 years

# Agenda

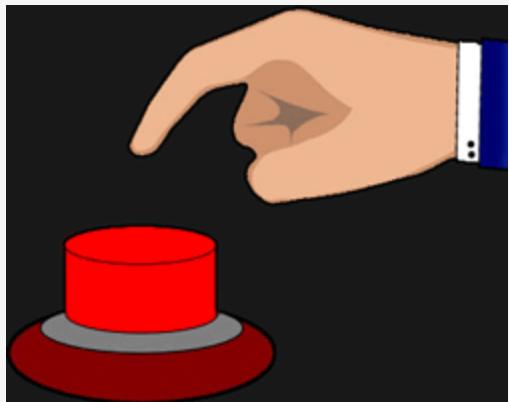
- Boring Bits
- Lessons Learnt
- Performance Engineering
- Wrap up

Quick Survey 

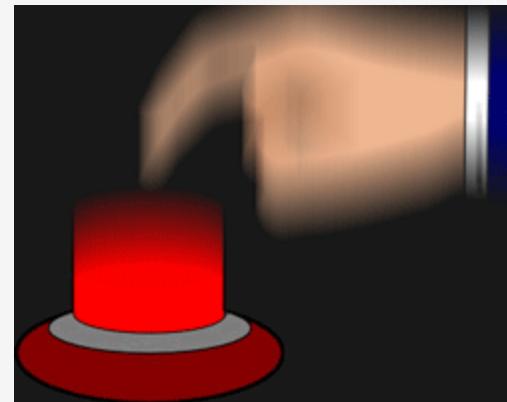
...a testing practice performed to determine how a system performs in terms of responsiveness and stability...

(source: wikipedia)

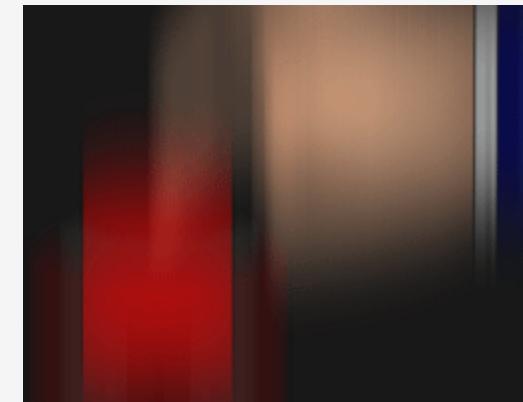
So like?



fast



faster



fastest

A white rabbit is sitting on a red and green checkered cloth. The rabbit is facing towards the right of the frame. Its fur is white, and it has dark eyes and ears. The background is slightly blurred.

What can we measure?

# What can we measure?

## Client

Latency/Response Time - Round Trip

Throughput

## Server

Latency/Response Time - Processing Time

Throughput

Availability

Server metrics - CPU, memory, disk I/O, network I/O...

Connection pooling

Cache hit/miss ratios  
and more...

# What flavours does it come in?

- Load testing
- Stress testing
- Soak testing
- Spike testing
- Breakpoint testing
- Configuration testing
- Isolation testing
- Internet testing

---

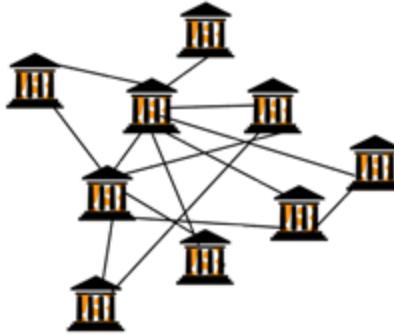
additional reading

The Art of Application Performance Testing - Ian Molyneaux





CENTRALIZED



DECENTRALIZED

EXCHANGE CONTROLS FUNDS

USER CONTROLS FUNDS

NOT ANONYMOUS

ANONYMOUS

HACKS & SERVER DOWNTIME

NO HACKS & SERVER DOWNTIME

source: [bitcoinwiki](#)

# API

## User

User Info

Accounts

Deposits

Withdrawals

## Trade

### Orders

[Place a new order](#)

Cancel an order

Cancel all orders

List Orders

Get V1 Historical Orders List

Recent Orders

Get an order

Fills

## Market Data

Symbols & Ticker

Order Book

Histories

Currencies

## Websocket Feed

### Public Channels

#### Symbol Ticker

All Symbols Ticker

Symbol Snapshot

Market Snapshot

Level-2 Market Data

Match Execution Data

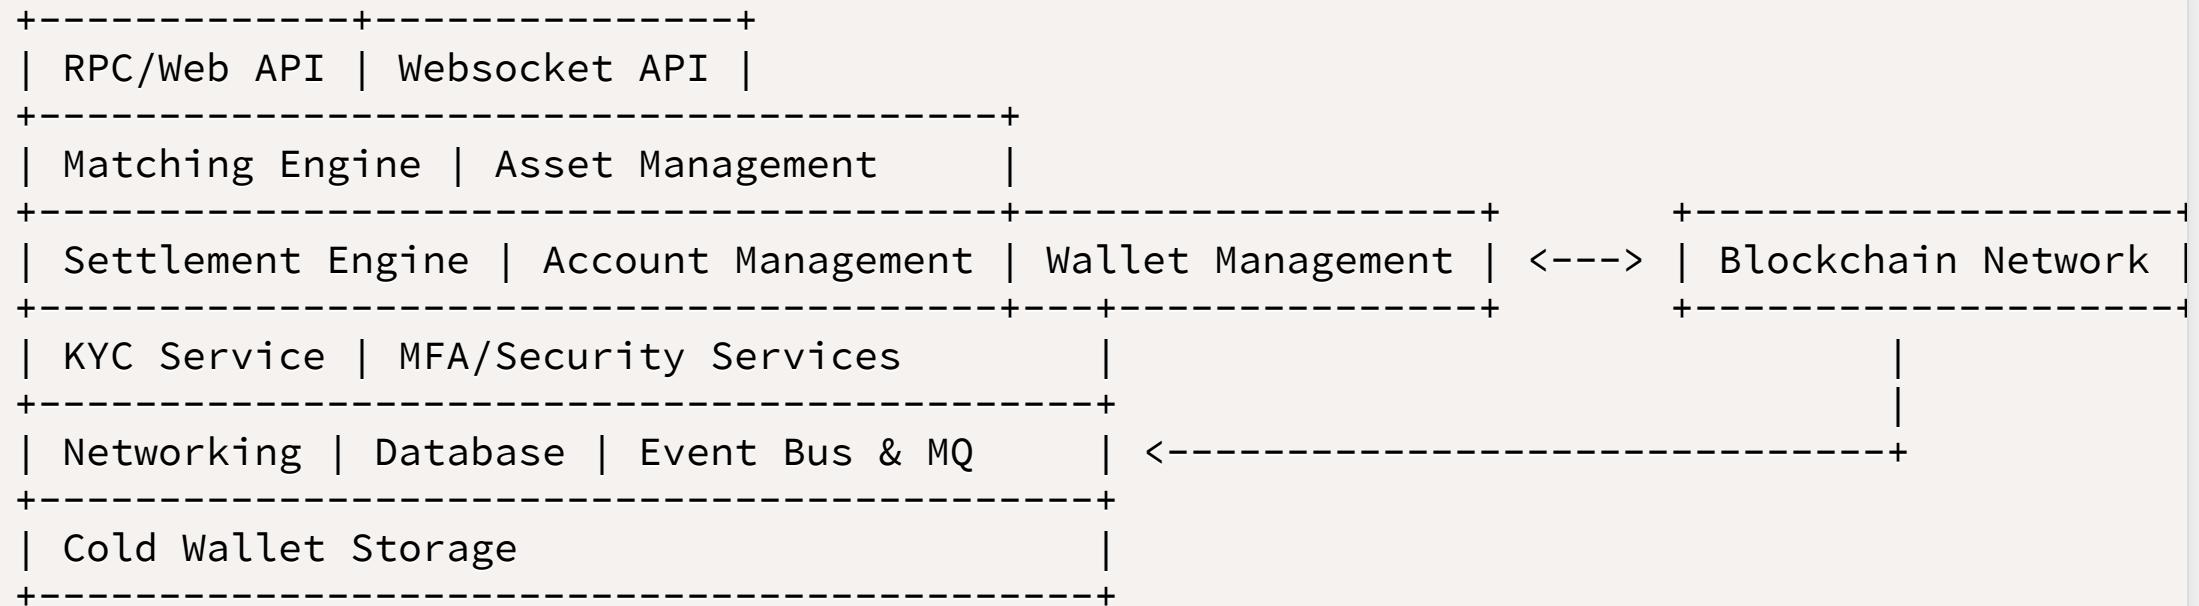
### Others

Time

Full MatchEngine Data(Level 3)

Private Channels

# Typical Architecture



additional reading: [How do cryptocurrency exchanges work? and what technologies are driving disruption - Naveen Saraswat](#)

# Let the show begin!

- Centralised crypto exchange API
- Order submission/execution
- Gatling



# Gatling

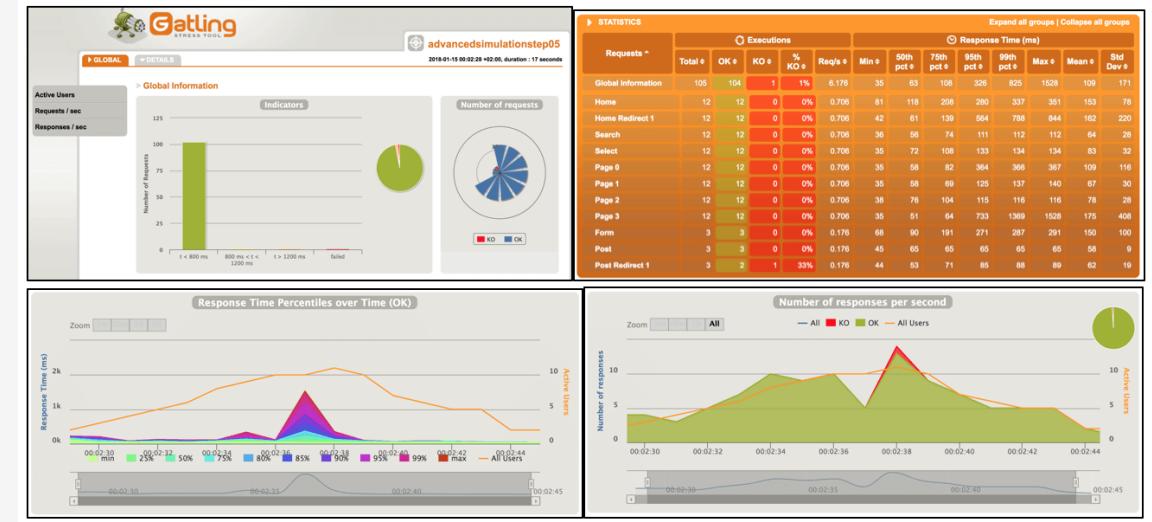
- non-blocking/asynchronous stack (scala, akka, netty)
- not just a load runner, can be scripted using DSL
- built-in assertion
- good calculation and statistics
- nice reporting
- comes with a recorder
- open-source\*
- can run in a distributed fashion and feed into other performance platforms

```
package computerdatabase // 1
import io.gatling.core.Predef._ // 2
import io.gatling.http.Predef_
import scala.concurrent.duration._

class BasicSimulation extends Simulation { // 3

    val httpProtocol = http // 4
        .baseUrl("http://computer-database.gatling.io") // 5
        .acceptHeader("text/html,application/xhtml+xml,application/xml;q=0.9,*/*;q=0.8") // 6
        .doNotTrackHeader("1")
        .acceptLanguageHeader("en-US,en;q=0.5")
        .acceptEncodingHeader("gzip, deflate")
        .userAgentHeader("Mozilla/5.0 (Windows NT 5.1; rv:31.0) Gecko/20100101 Firefox/31.0")

    val scn = scenario("BasicSimulation") // 7
        .exec(http("request_1") // 8
            .get("/")) // 9
}
```





# Non-Functional Requirements

Expensive & Time Consuming

Baseline!

Without data you're just another person with an opinion

- W. Edwards Deming



Consider all the parameters



```
POST https://{{host}}/api/order
jwt: abcdefg1234567890
Content-Type: application/json
Accept: application/json
Accept-Charset: utf-8

{
    "symbol": "EthBtc",
    "side": "Buy",
    "order_type": "Limit",
    "time_in_force": "Gtc",
    "quantity": 2,
    "price": 10.3123,
    "new_order_resp_type": "Ack",
    "timestamp": 1562336244
}
```

# Parameters

```
val price = ??  
val quantity = ??  
  
val Side = Array("Buy", "Sell")  
val OrderTypes = Array("Market", "Limit", "StopLoss", "StopLossLimit", "TakeProfit", "TakeF  
val TimeInForce = Array("Gtc", "Ioc")  
val Symbols = Array("EthBtc", "EthLtc", "EthUsdt", "EthXrp")
```

```
exec(http("""POST /api/order STOPLOSS SELL ${symbol}""")  
    .post("/order")  
    .body(StringBody(  
        """{  
            "symbol": "${symbol}",  
            "side": "${side}",  
            "order_type": "StopLoss",  
            "time_in_force": "Gtc",  
            "quantity": ${quantity},  
            "stop_price": ${stopSellPrice},  
            "new_order_resp_type": "Ack",  
            "timestamp": ${timestamp}  
        }""")).asJson  
    .headers(PostHeaders)  
    .header("jwt", """${jwt}""")  
    .check(status.is(200))  
)
```



```
val orderParams: Iterator[Map[String, Any]] = Iterator.continually(  
  Map(  
    "quantity" -> Random.nextDouble() * 100,  
    "price" -> Random.nextDouble() * 10,  
    "symbol" -> Symbols(Random.nextInt(Symbols.length)),  
    "timestamp" -> Instant.now.getEpochSecond,  
  )  
)
```



```
val orderParams: Iterator[Map[String, Any]] = Iterator.continually(  
    elem = {  
        val symbol = Random.nextInt(Symbols.length)  
        val marketPrice = getMarketPrice(symbol)  
        Map(  
            "quantity" -> Random.nextDouble() * 100,  
            "price" -> marketPrice,  
            "sellPrice" -> marketPrice * (1 + (Random.nextInt(5) / 1000)),  
            "buyPrice" -> marketPrice * (1 - (Random.nextInt(5) / 1000)),  
            "symbol" -> Symbols(symbol),  
            "timestamp" -> Instant.now.getEpochSecond,  
        )  
    }  
)
```



## Add some asserts!

- response code
- check orders matched and reconcile with user balance



# Trading Bots

additional functional tests while under load



Know the limits

- Rate limiting
- Quota
- DDOS protection

Fast & Slow Users





Trust, but Verify (Developers)

Developer's performance intuitions are often wrong

Myself included

# Things Devs might say

- "Get a bigger instance with more ram/cpu"
- "Just add more instances"
- "It must be the database"
- "It's slow but it hasn't crashed yet"
- "Must be an environment issue"
- "Your test is wrong"

▶ STATISTICS														<a href="#">Expand all groups</a>   <a href="#">Collapse all groups</a>		
Requests ^	⌚ Executions						⌚ Response Time (ms)									
	Total ↴	OK ↴	KO ↴	% KO ↴	Req/s ↴	Min ↴	50th pct ↴	75th pct ↴	95th pct ↴	99th pct ↴	Max ↴	Mean ↴	Std Dev ↴			
Global Information	7167293	7167293	0	0%	3643.769	1	12	28	71	134	7273	24	65			
POST /ap...Y EthXrp	452846	452846	0	0%	230.222	1	9	25	67	133	7261	23	67			
POST /ap...L EthXrp	443799	443799	0	0%	225.622	1	13	31	73	134	7227	26	66			
POST /ap...L EthBtc	450083	450083	0	0%	228.817	1	9	24	67	132	3062	22	64			
POST /ap...Y EthBtc	445321	445321	0	0%	226.396	1	14	32	74	136	3083	26	65			
POST /ap...Y EthLtc	445158	445158	0	0%	226.313	1	14	31	73	134	3058	26	65			
POST /ap...L EthLtc	451691	451691	0	0%	229.634	1	9	24	67	132	3081	22	64			
POST /ap...L EthLtc	444752	444752	0	0%	226.107	1	13	31	73	135	3061	26	66			
POST /ap... EthUsdt	445641	445641	0	0%	226.559	1	14	32	75	136	7190	27	65			
POST /ap...Y EthXrp	445064	445064	0	0%	226.265	1	14	31	73	133	7273	26	64			
POST /ap...L EthBtc	443901	443901	0	0%	225.674	1	13	31	74	135	3022	26	64			
POST /ap...Y EthLtc	451208	451208	0	0%	229.389	1	9	25	67	132	2225	22	64			
POST /ap... EthUsdt	444272	444272	0	0%	225.863	1	14	32	75	136	3393	27	65			
POST /ap...Y EthBtc	450960	450960	0	0%	229.263	1	9	25	67	133	7123	23	67			
POST /ap... EthUsdt	450506	450506	0	0%	229.032	1	9	24	67	132	3085	22	65			
POST /ap...L EthXrp	450951	450951	0	0%	229.258	1	9	24	67	132	3082	22	64			
POST /ap... EthUsdt	451140	451140	0	0%	229.354	1	9	24	67	133	3107	23	65			



```
1% tile: 974.0 (ns)
5% tile: 1075.0 (ns)
10% tile: 2292.0 (ns)
25% tile: 2695.0 (ns)
50% tile: 3671.0 (ns)
75% tile: 10440.0 (ns)
90% tile: 10091923.79999999 (ns)
99% tile: 68835579.60000025 (ns)
```

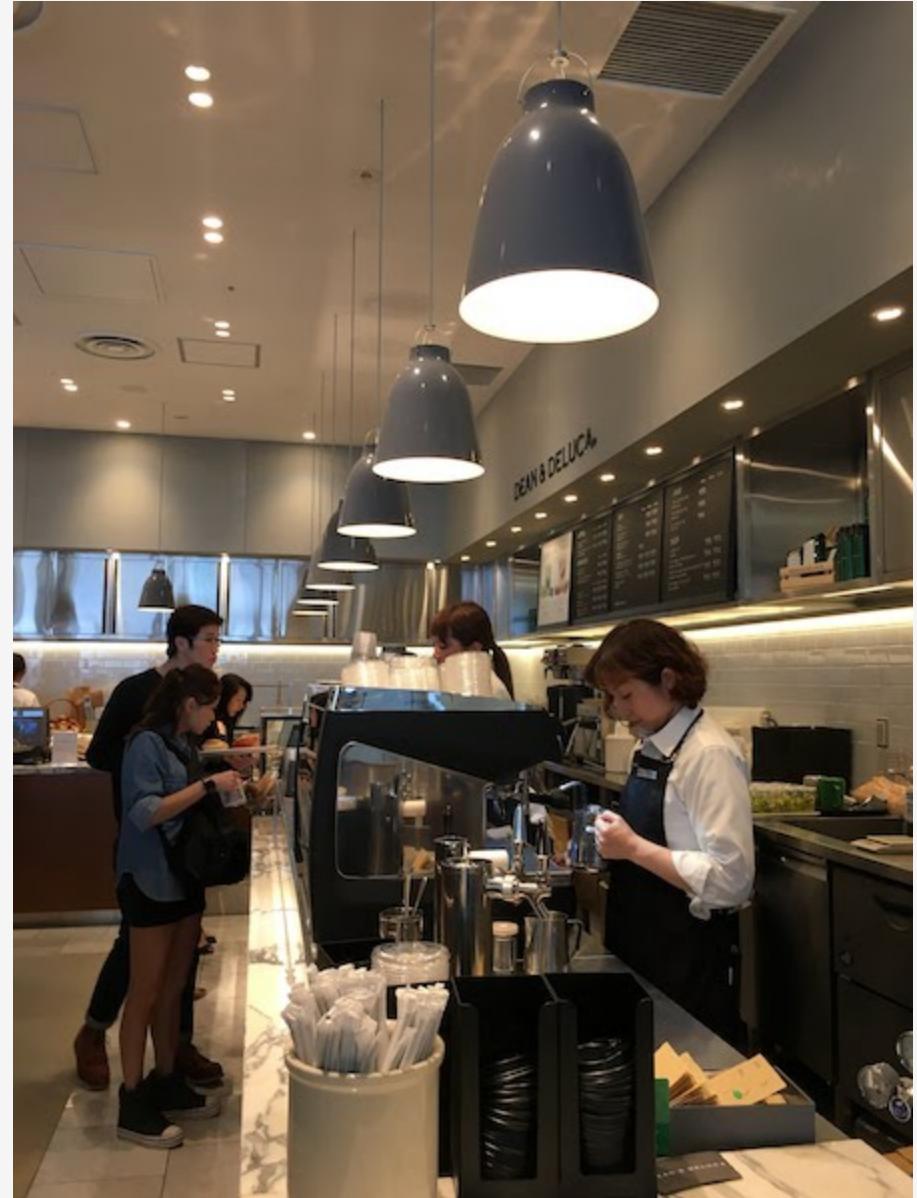
# Things Devs might do

```
let start = Instant::now();
submit_order();
let duration = start.elapsed();

println!("Time elapsed in submit_order() is: {:?}", duration);
```

# Coordinated Omission Problem

additional reading: ["How NOT to Measure Latency" by Gil Tene](#)



Application Performance Monitoring Metrics

+

Log Aggregation w/ Correlation IDs

=

Observability



Collect Data



Find the bottleneck



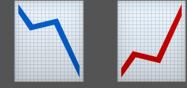
Fix it



Repeat

Without data you're just another person with an opinion.

- W. Edwards Deming



Statistics Lie

▶ STATISTICS														<a href="#">Expand all groups</a>   <a href="#">Collapse all groups</a>		
Requests ^	⌚ Executions						⌚ Response Time (ms)									
	Total ↴	OK ↴	KO ↴	% KO ↴	Req/s ↴	Min ↴	50th pct ↴	75th pct ↴	95th pct ↴	99th pct ↴	Max ↴	Mean ↴	Std Dev ↴			
Global Information	7167293	7167293	0	0%	3643.769	1	12	28	71	134	7273	24	65			
POST /ap...Y EthXrp	452846	452846	0	0%	230.222	1	9	25	67	133	7261	23	67			
POST /ap...L EthXrp	443799	443799	0	0%	225.622	1	13	31	73	134	7227	26	66			
POST /ap...L EthBtc	450083	450083	0	0%	228.817	1	9	24	67	132	3062	22	64			
POST /ap...Y EthBtc	445321	445321	0	0%	226.396	1	14	32	74	136	3083	26	65			
POST /ap...Y EthLtc	445158	445158	0	0%	226.313	1	14	31	73	134	3058	26	65			
POST /ap...L EthLtc	451691	451691	0	0%	229.634	1	9	24	67	132	3081	22	64			
POST /ap...L EthLtc	444752	444752	0	0%	226.107	1	13	31	73	135	3061	26	66			
POST /ap... EthUsdt	445641	445641	0	0%	226.559	1	14	32	75	136	7190	27	65			
POST /ap...Y EthXrp	445064	445064	0	0%	226.265	1	14	31	73	133	7273	26	64			
POST /ap...L EthBtc	443901	443901	0	0%	225.674	1	13	31	74	135	3022	26	64			
POST /ap...Y EthLtc	451208	451208	0	0%	229.389	1	9	25	67	132	2225	22	64			
POST /ap... EthUsdt	444272	444272	0	0%	225.863	1	14	32	75	136	3393	27	65			
POST /ap...Y EthBtc	450960	450960	0	0%	229.263	1	9	25	67	133	7123	23	67			
POST /ap... EthUsdt	450506	450506	0	0%	229.032	1	9	24	67	132	3085	22	65			
POST /ap...L EthXrp	450951	450951	0	0%	229.258	1	9	24	67	132	3082	22	64			
POST /ap... EthUsdt	451140	451140	0	0%	229.354	1	9	24	67	133	3107	23	65			

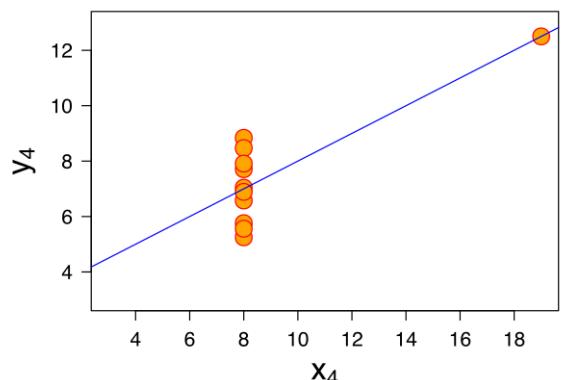
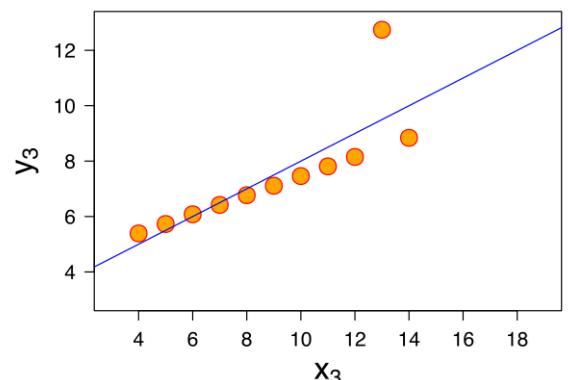
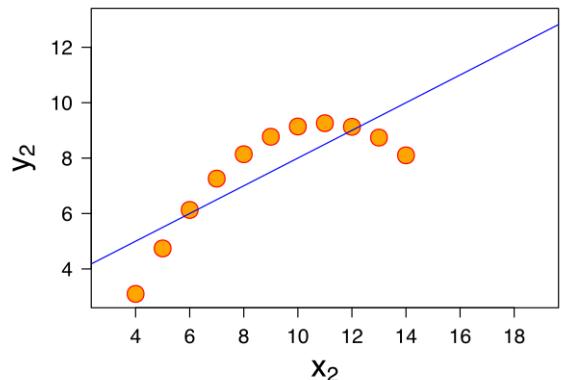
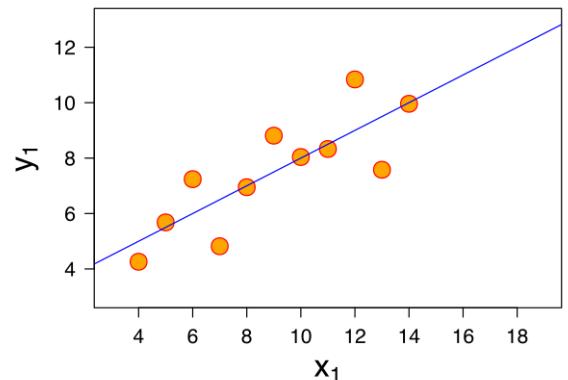
Ignore

✗ Mean

✗ Median

✗ Standard Deviation

# Anscombe's Quartet



	I	II	III	IV	
x	x	x	x	x	
10	8,04	10	9,14	10	7,46
8	6,95	8	8,14	8	6,77
13	7,58	13	8,74	13	12,74
9	8,81	9	8,77	9	7,11
11	8,33	11	9,26	11	7,81
14	9,96	14	8,1	14	8,84
6	7,24	6	6,13	6	6,08
4	4,26	4	3,1	4	5,39
12	10,84	12	9,13	12	8,15
7	4,82	7	7,26	7	6,42
5	5,68	5	4,74	5	5,73
SUM	99,00	82,51	99,00	82,51	99,00
AVG	9,00	7,50	9,00	7,50	9,00
STDEV	3,32	2,03	3,32	2,03	3,32

Look at

- Max
- Percentiles

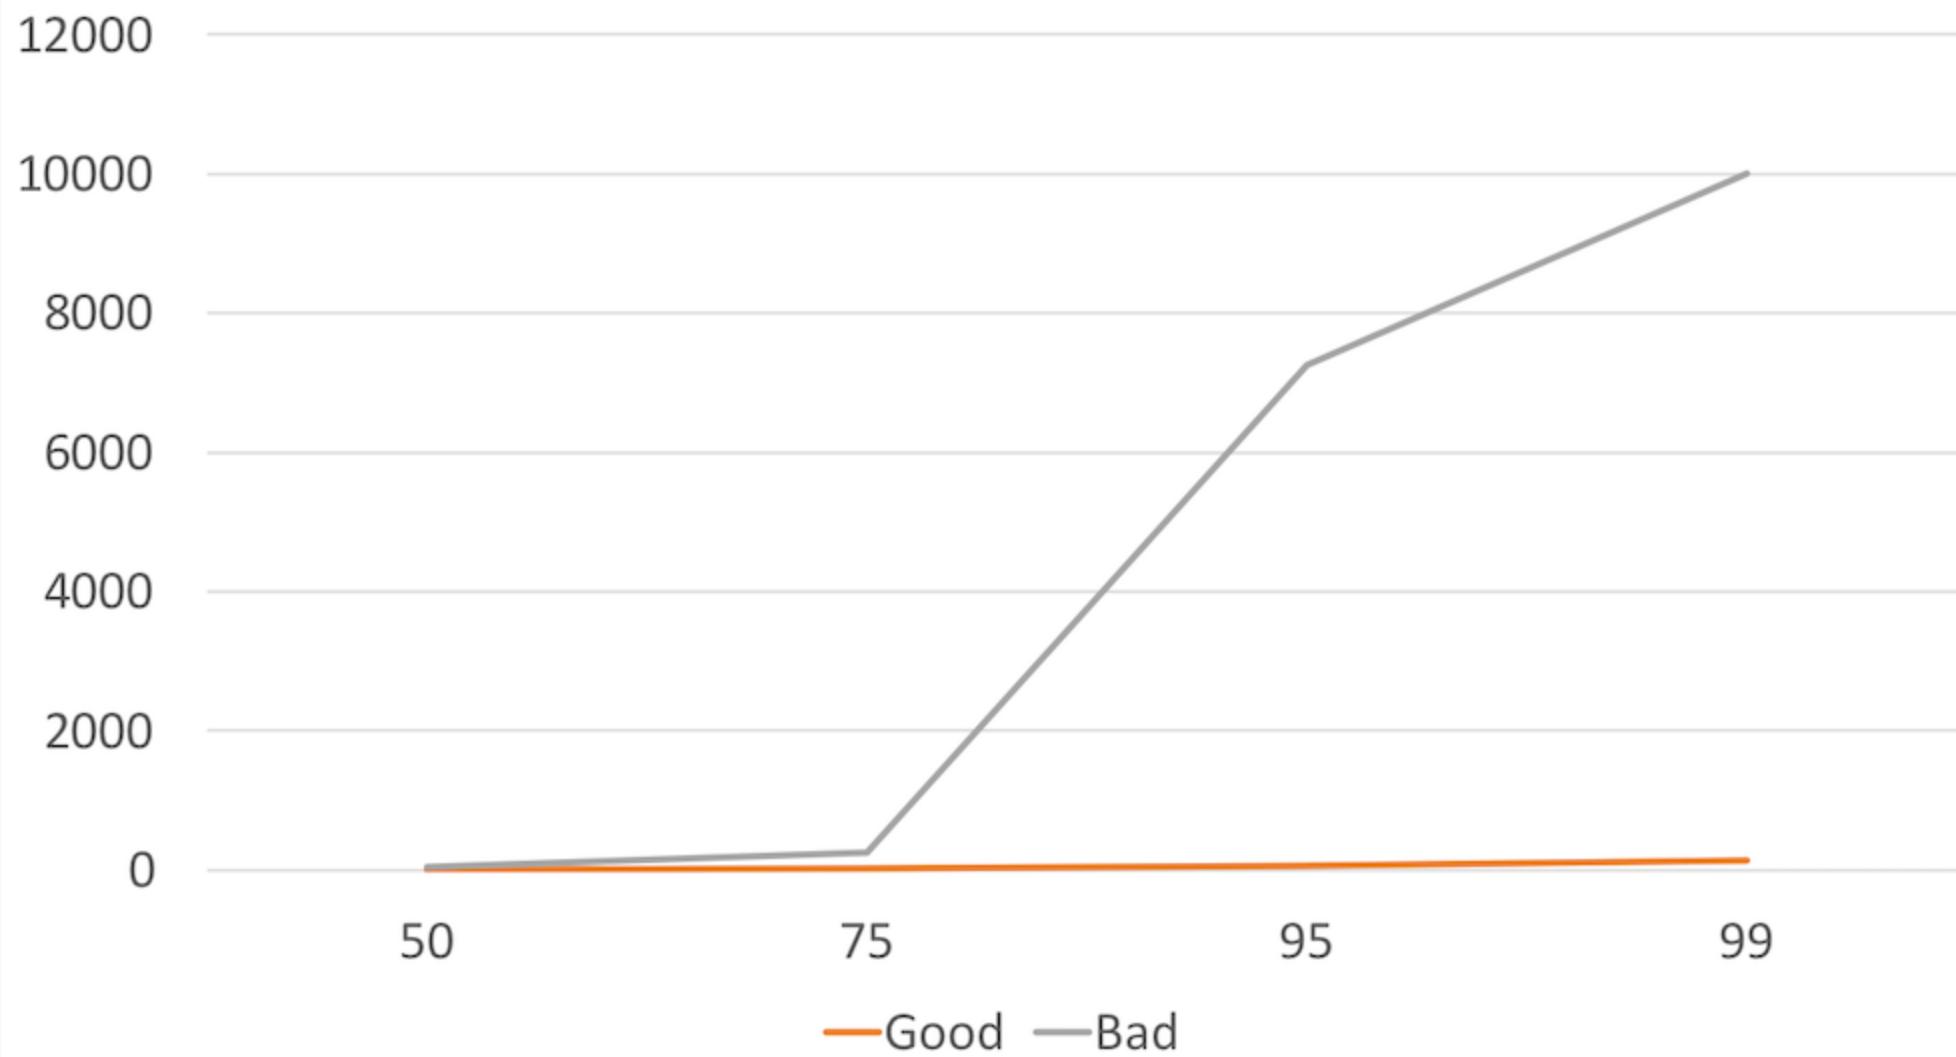
▶ STATISTICS														<a href="#">Expand all groups</a>   <a href="#">Collapse all groups</a>		
Requests ^	⌚ Executions						⌚ Response Time (ms)									
	Total ↴	OK ↴	KO ↴	% KO ↴	Req/s ↴	Min ↴	50th pct ↴	75th pct ↴	95th pct ↴	99th pct ↴	Max ↴	Mean ↴	Std Dev ↴			
Global Information	7167293	7167293	0	0%	3643.769	1	12	28	71	134	7273	24	65			
POST /ap...Y EthXrp	452846	452846	0	0%	230.222	1	9	25	67	133	7261	23	67			
POST /ap...L EthXrp	443799	443799	0	0%	225.622	1	13	31	73	134	7227	26	66			
POST /ap...L EthBtc	450083	450083	0	0%	228.817	1	9	24	67	132	3062	22	64			
POST /ap...Y EthBtc	445321	445321	0	0%	226.396	1	14	32	74	136	3083	26	65			
POST /ap...Y EthLtc	445158	445158	0	0%	226.313	1	14	31	73	134	3058	26	65			
POST /ap...L EthLtc	451691	451691	0	0%	229.634	1	9	24	67	132	3081	22	64			
POST /ap...L EthLtc	444752	444752	0	0%	226.107	1	13	31	73	135	3061	26	66			
POST /ap... EthUsdt	445641	445641	0	0%	226.559	1	14	32	75	136	7190	27	65			
POST /ap...Y EthXrp	445064	445064	0	0%	226.265	1	14	31	73	133	7273	26	64			
POST /ap...L EthBtc	443901	443901	0	0%	225.674	1	13	31	74	135	3022	26	64			
POST /ap...Y EthLtc	451208	451208	0	0%	229.389	1	9	25	67	132	2225	22	64			
POST /ap... EthUsdt	444272	444272	0	0%	225.863	1	14	32	75	136	3393	27	65			
POST /ap...Y EthBtc	450960	450960	0	0%	229.263	1	9	25	67	133	7123	23	67			
POST /ap... EthUsdt	450506	450506	0	0%	229.032	1	9	24	67	132	3085	22	65			
POST /ap...L EthXrp	450951	450951	0	0%	229.258	1	9	24	67	132	3082	22	64			
POST /ap... EthUsdt	451140	451140	0	0%	229.354	1	9	24	67	133	3107	23	65			

## ▶ STATISTICS

[Expand all groups](#) | [Collapse all groups](#)

Requests	🕒 Executions					⌚ Response Time (ms)							
	Total	OK	KO	% KO	Req/s	Min	50th pct	75th pct	95th pct	99th pct	Max	Mean	Std Dev
Global Information	2706735	2615596	91139	3%	1372.584	1	45	264	7262	10003	60076	1089	4294
POST /ap... EthUsdt	338833	327372	11461	3%	171.822	1	47	268	7280	10003	60035	1099	4331
POST /ap...Y EthBtc	338607	327161	11446	3%	171.707	1	44	262	7255	10003	60076	1086	4292
POST /ap...L EthXrp	338458	326961	11497	3%	171.632	1	46	269	7285	10003	60021	1100	4338
POST /ap...Y EthXrp	338447	327118	11329	3%	171.626	1	46	265	7256	10003	60037	1093	4313
POST /ap...L EthBtc	338319	326995	11324	3%	171.561	1	44	264	7253	10003	60036	1084	4271
POST /ap... EthUsdt	338157	326779	11378	3%	171.479	1	47	264	7242	10003	60028	1084	4289
POST /ap...L EthLtc	337958	326695	11263	3%	171.378	1	44	263	7241	10003	60025	1075	4222
POST /ap...Y EthLtc	337956	326515	11441	3%	171.377	1	44	262	7285	10003	60023	1092	4293

## Percentile Plot





# Scaling your performance test

Focus on the actual test

Scaling your test comes later

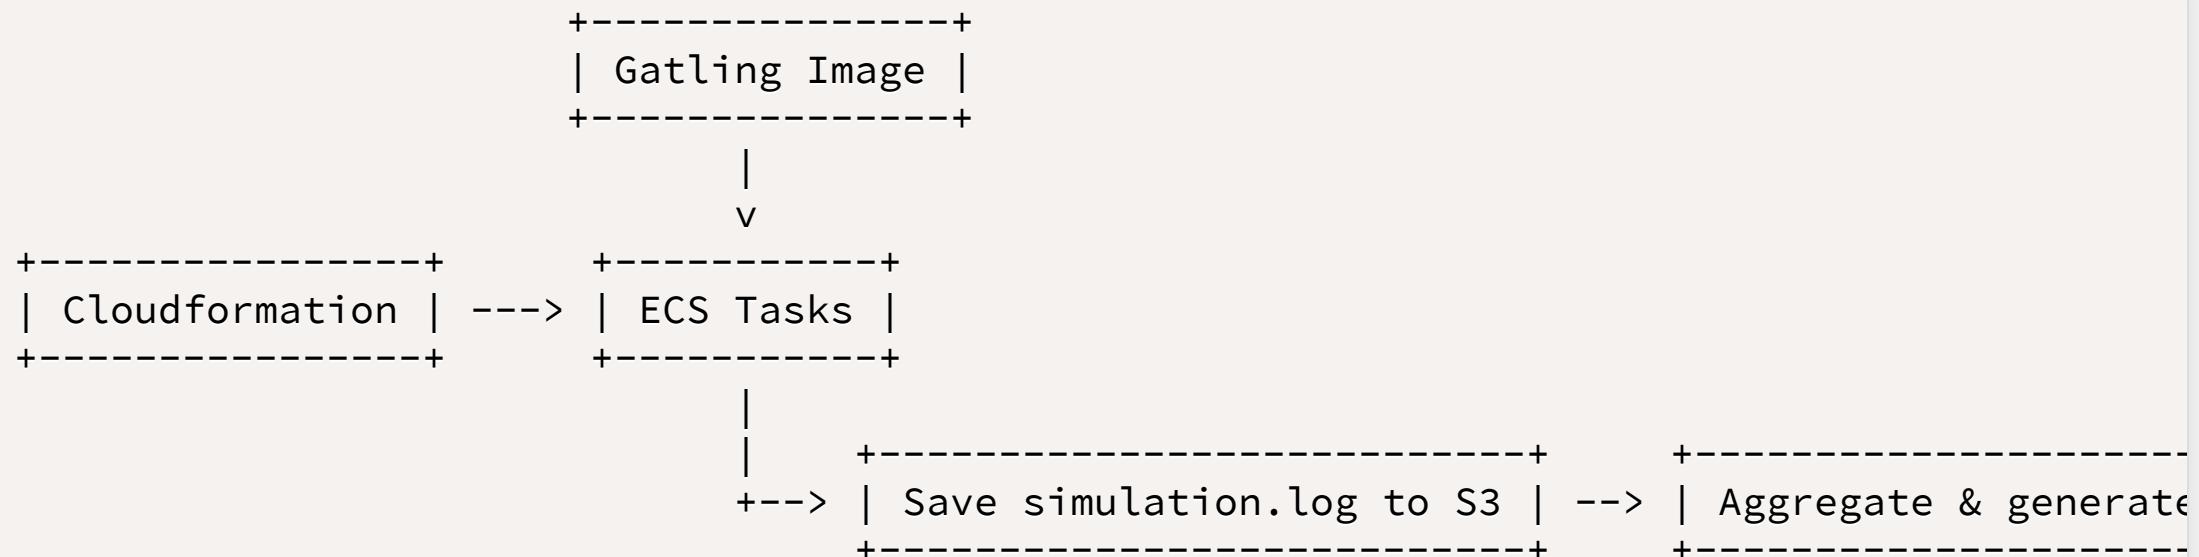
# Paid Performance Testing Platform

Gatling Frontline

BlazeMeter

Flood.io

# DIY - (The Hard way)





# Performance Engineering

Quickfire Edition

## What about Microservices/Serverless?

- Watch for timeouts
- Perform spike tests to determine time taken to scale

## What about GraphQL?

- It's possible, but more complicated
- Different combination of fields and fragments
- Make sure you have sensible request tracing

# Many birds, one stone

Reusable Tests. Framework selection is critical!

E2E test in Dev/Test Deployment

```
./gradlew gatlingRun -DbaseUrl="http://dev.env:80/api" -DnumberOfUser=1 -DrunDurationSecs=3
```

Perf Test

```
./gradlew gatlingRun -DbaseUrl="http://perf.env:80/api" -DnumberOfUser=2000 -DrunDurationSecs=30
```

Smoke Test in Production

```
./gradlew gatlingRun -DbaseUrl="http://prod.env:80/api" -DnumberOfUser=1 -DrunDurationSecs=3
```

## CI/CD

- Build a small subset of your performance testing suite as part of your pipeline
- Monitor the build time and capture performance metrics
- Fail or Add alert for any executions that over n %
- Run perf early in the SDLC and as often as possible

# It doesn't take a lot to cause an outage

A regular expression that backtracked enormously and exhausted CPU used for HTTP/HTTPS serv



Introduce performance profiling for all rules to the test suite. (ETA: July 19)

Test in Production?

# Test in Production?

Yes, only if you can cleanup the data

Otherwise, test in a separate environment with configuration similar to prod

# Takeaways

- Performance Test != Non-Functional Testing
- Know your parameters and limits upfront
- Test with the end in mind
- Forgo any assumptions and verify using performance metrics
- Stats Lie!

# Awesome Performance

<https://github.com/hbish/awesome-performance>



Thanks

[@hbish](https://twitter.com/hbish) | [/in/benshi](https://www.linkedin.com/in/benshi) | [hbish.com](https://hbish.com)